

A Distributed Algorithm for Minimum-Weight Spanning Trees

by

R. G. Gallager, P.A. Humblet, and P. M. Spira

ACM, Transactions on Programming Language and systems, 1983

presented by
Hanan Shpungin

Outline


- Introduction
- The idea of Distributed MST
- The algorithm

Outline

- **Introduction**
- The idea of Distributed MST
- The algorithm

Introduction


- Problem
 - A graph $G = (V, E)$
 - Every edge has a weight $\forall e \in E, w(e) \in \mathbb{R}$



Introduction

- Solution


- A spanning tree $T = (V, E')$
- So that the sum $\sum_{e \in E'} w(e)$ is minimized



Introduction

- Solution

- A spanning tree $T = (V, E')$
- So that the sum $\sum_{e \in E'} w(e)$ is minimized




Introduction

- Definition – MST *fragment*

- A connected sub-tree of MST


Example of possible fragments:



Introduction

- MST Property 1

Given a fragment of an MST, let e be a minimum-weight outgoing edge of the fragment. Then joining e and its adjacent non-fragment node to the fragment yields another fragment of an MST.




Introduction

- MST Property 1

Given a fragment of an MST, let e be a minimum-weight outgoing edge of the fragment. Then joining e and its adjacent non-fragment node to the fragment yields another fragment of an MST.

- *Proof:*




Introduction

- MST Property 1

Given a fragment of an MST, let e be a minimum-weight outgoing edge of the fragment. Then joining e and its adjacent non-fragment node to the fragment yields another fragment of an MST.

- *Proof:*

Suppose e is not in MST, but some e' instead.



Introduction


- **MST Property 1**

Given a fragment of an MST, let e be a minimum-weight outgoing edge of the fragment. Then joining e and its adjacent non-fragment node to the fragment yields another fragment of an MST.

- **Proof:**

Suppose e is not in MST, but some e' instead.

MST with e forms a cycle.



Introduction

- **MST Property 1**


Given a fragment of an MST, let e be a minimum-weight outgoing edge of the fragment. Then joining e and its adjacent non-fragment node to the fragment yields another fragment of an MST.

- **Proof:**

Suppose e is not in MST, but some e' instead.

MST with e forms a cycle.


We obtain a cheaper MST with e instead of e' .



Introduction

- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.




Introduction

- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*





Introduction

- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*

Suppose existence of two MSTs.



Introduction



- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*

Suppose existence of two MSTs.

Let e be the minimal-weight edge not in both MSTs (wlog e in T).



Introduction

- MST Property 2



If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*

Suppose existence of two MSTs.

Let e be the minimal-weight edge not in both MSTs (wlog e in T).

T' with e has a cycle




Introduction


- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*



*Suppose existence of two MSTs.
Let e be the minimal-weight edge
not in both MSTs (wlog e in T).
 T' with e has a cycle
At least cycle edge e' is not in T .*




Introduction

- MST Property 2

If all the edges of a connected graph have different weights, then the MST is unique.

- *Proof:*




Suppose existence of two MSTs.

Let e be the minimal-weight edge not in both MSTs (wlog e in T).

T' with e has a cycle


At least cycle edge e' is not in T .

Since $w(e) < w(e')$ we conclude that T' with e and without e' is a smaller MST than T' .




Introduction

- Idea of MST based on properties 1 & 2
 - Start with fragments of one node.




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)



Introduction

- Idea of MST based on properties 1 & 2
 - Enlarge fragments in any order (property 1)
 - Combine fragments with a common node (property 2)




Outline

- Introduction
- **The idea of Distributed MST**
- The algorithm


The idea of Distributed MST

- Fragments
 - Every node starts as a single fragment.




The idea of Distributed MST

- Fragments
 - Each fragment finds its minimum outgoing edge.



The idea of Distributed MST

- Fragments
 - Each fragment finds its minimum outgoing edge.
 - Then it tries to combine with the adjacent fragment.




The idea of Distributed MST

- Levels
 - Every fragment has an associated level that has impact on combining fragments.
 - A fragment with a single node is defined to be at level 0.

The idea of Distributed MST


- Levels
 - The combination of two fragments depends on the levels of fragments.



The idea of Distributed MST

- Levels
 - The combination of two fragments depends on the levels of fragments.

If a fragment F wishes to connect to a fragment F' and $L < L'$ then:




The idea of Distributed MST

- Levels
 - The combination of two fragments depends on the levels of fragments.

If a fragment F wishes to connect to a fragment F' and $L < L'$ then:


F is absorbed in F' and the resulting fragment is at level L' .



The idea of Distributed MST

- Levels
 - The combination of two fragments depends on the levels of fragments.

If fragments F and F' have the same minimum outgoing edge and $L = L'$ then:




The idea of Distributed MST

- Levels
 - The combination of two fragments depends on the levels of fragments.

If fragments F and F' have the same minimum outgoing edge and $L = L'$ then:


The fragments combine into a new fragment F'' at level $L'' = L + 1$.



The idea of Distributed MST

- Levels
 - The identity of a fragment is the weight of its *core*.

If fragments F and F' with same level were combined, the combining edge is called the core of the new segment.



The idea of Distributed MST

- State
 - Each node has a state

Sleeping - initial state

Find - during fragment's search for a minimal outgoing edge


Found - otherwise (when a minimal outgoing edge was found)

Outline

- Introduction
- The idea of Distributed MST
- **The algorithm**

The Algorithm


- Fragment minimum outgoing edge discovery
 - Special case of zero-level fragment (*Sleeping*).



The Algorithm

- Fragment minimum outgoing edge discovery
 - Special case of zero-level fragment (*Sleeping*).

*When a node awakes from the state **Sleeping**, it finds a minimum edge connected.*




The Algorithm

- Minimum outgoing edge discovery
 - Special case of zero-level fragment (*Sleeping*).

*When a node awakes from the state **Sleeping**, it finds a minimum edge connected.*

*Marks it as a **branch** of MST and sends a **Connect** message over this edge.*




The Algorithm

- Minimum outgoing edge discovery
 - Special case of zero-level fragment (*Sleeping*).

*When a node awakes from a state **Sleeping**, it finds a minimum edge connected.*


*Marks it as a **branch** of MST and sends a **Connect** message over this edge.*

*Goes into a **Found** state.*



The Algorithm

- Minimum outgoing edge discovery
 - Take a fragment at level L that was just combined out of two level $L-1$ fragments.




The Algorithm

- Minimum outgoing edge discovery
 - Take a fragment at level L that was just combined out of two level $L-1$ fragments.

The weight of the core is the identity of the fragment.


It acts as a root of a fragment tree.



The Algorithm

- Minimum outgoing edge discovery
 - Take a fragment at level L that was just combined out of two level L-1 fragments.

*Nodes adjacent to core send an **Initiate** message to the borders.*




The Algorithm

- Minimum outgoing edge discovery
 - Take a fragment at level L that was just combined out of two level L-1 fragments.

*Nodes adjacent to core send an **Initiate** message to the borders.*

Relayed by the intermediate nodes in the fragment.




The Algorithm

- Minimum outgoing edge discovery
 - Take a fragment at level L that was just combined out of two level L-1 fragments.

*Nodes adjacent to core send an **Initiate** message to the borders.*

Relayed by the intermediate nodes in the fragment.


*Puts the nodes in the **Find** state.*



The Algorithm

- Minimum outgoing edge discovery
 - Edge classification/


Basic - yet to be classified, can be inside fragment or outgoing edges



The Algorithm

- Minimum outgoing edge discovery
 - Edge classification.


Rejected – an inside fragment edge



The Algorithm

- Minimum outgoing edge discovery
 - Edge classification.


Branch – an MST edge



The Algorithm

- Minimum outgoing edge discovery
 - On receiving the *Initiate* message a node tries to find a minimum outgoing edge.

*Sends a **Test** message on
Basic edges (minimal first)*




The Algorithm

- Minimum outgoing edge discovery
 - On receiving the *Test* message.

*In case of same identity:
send a **Reject** message, the
edge is **Rejected**.*


*In case **Test** was sent in
both directions, the edge is
Rejected automatically
without a **Reject** message.*



The Algorithm

- Minimum outgoing edge discovery
 - On receiving the *Test* message.


*In case of a self lower level:
Delay the response until the
identity rises sufficiently.*



The Algorithm

- Minimum outgoing edge discovery
 - On receiving the *Test* message.

*In case of a self higher level:
Send an **Accept** message
The edge is accepted as a
candidate.*




The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.

*Nodes send **Report** messages along the branches of the MST.*

If no outgoing edge was found the algorithm is complete.




The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.

*Nodes send **Report** messages along the branches of the MST.*

If no outgoing edge was found the algorithm is complete.


*After sending they go into **Found** mode.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.


*Every leaf sends the **Report** when resolved its outgoing edge.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.


*Every leaf sends the **Report** when resolved its outgoing edge.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.


*Every interior sends the **Report** when resolved its outgoing and all its children sent theirs.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.


*Every interior sends the **Report** when resolved its outgoing and all its children sent theirs.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.


*Every node remembers the branch to the minimal outgoing edge of its sub-tree, denoted **best-edge**.*



The Algorithm

- Minimum outgoing edge discovery
 - Agreeing on the minimal outgoing edge.

*The core adjacent nodes exchange **Reports** and decide on the minimal outgoing edge.*




The Algorithm

- Combining segments
 - Changing core.

*When decided a **Change-core** message is sent over branches to the minimal outgoing edge.*

*The tree **branches** point to the new core.*




The Algorithm

- Combining segments
 - Changing core

*When decided a **Change-core** message is sent over branches to the minimal outgoing edge.*


*The tree **branches** point to the new core.*

*Finally a **Connect** message is sent over the minimal edge.*



The Algorithm

- Final notes
 - Connecting same level fragments.




The Algorithm

- Final notes
 - Connecting same level fragments.

*Both core adjacent nodes send a **Connect** message, which causes the level to be increased.*

*As a result, core is changed and $L=1$ new **Initiate** messages are sent.*



The Algorithm

- Final notes
 - Connecting lower level fragments.

*When lower level fragment F' at node n' joins some fragment F at node n before n sent its **Report**.*

*We can send n' an **Initiate** message with the **Find** state so it joins the search.*

The Algorithm

- Final notes
 - Connecting lower level fragments.

*When lower level fragment F' at node n' joins some fragment F at node n after n sent its **Report**.*

*It means that n already found a lower edge and therefore we can send n' an **Initiate** message with the **Found** state so it doesn't join the search.*

The Algorithm

- Final notes
 - Forwarding the *Initiate* message at level L .

*When forwarding an **Initiate** message to the leafs, it is also forwarded to any pending fragments at level $L-1$, as they might be delayed with response.*

The Algorithm

- Final notes
 - Upper bound on fragment levels.

Level $L+1$ contains at least 2 fragments at level L .

Level L contains at least 2^L nodes.

$\log_2 N$ is an upper bound on fragment levels.

The Algorithm

- Proof outline
 - Correct MST build-up

*The **Connect** message is sent on the minimal outgoing edge only.*

As a result of properties 1 & 2 we should obtain an MST.

The Algorithm

- Proof outline
 - No deadlocks

Assume there is a deadlock.

*Choose a fragment from the lowest level set
and with minimal outgoing edge in that set.*

*Its **Test/Connect** message surely will be replied.*

There will always be a working fragment.

The Algorithm

- Complexity
 - Communication

*At most E **Reject** messages (with corresponding E **Test** messages, because each edge can be rejected only once.*

The Algorithm

- Complexity
 - Communication

*At every but the zero or last levels, each node can accept up to 1 **Initiate**, **Accept** messages. It can transmit up to 1 **Test (successful)**, **Report**, **ChangeRoot**, **Connect**. Since the number of levels is bounded by $\log_2 N$ number of such messages is at most $5N(\log_2 N - 1)$.*

The Algorithm

- Complexity
 - Communication

*At level zero, each node receives at most one **Initiate** and transmits at most one **Connect**. At the last level a node can send at most one **Report** message, as a result at most $3N$ such messages.*

The Algorithm

- Complexity
 - Communication

As a result, the upper bound is: $5N \log_2 N + 2E$.

The Algorithm

- Complexity
 - Time (under assumption of initial awakening it is $5N \log_2 N$)

We prove by induction that one needs $5lN - 3N$ time units for every node to reach level l .

The Algorithm

- Complexity

- Time (under assumption of initial awakening it is $5N \log_2 N$)

*$l=1 \rightarrow$ Each node is awakened and sends a **Connect** message. By time $2N$ all nodes should be at level 1.*

The Algorithm

- Complexity

- Time (under assumption of initial awakening it is $5N \log_2 N$)

*Assume $l \rightarrow$ At level l , each node can send at most N **Test** messages which will be answered before time $5lN - N$. The propagation of the **Report** messages, **ChangeRoot**, **Connect**, and **Initiate** messages can take at most $3N$ units, so that by time $5(l + 1)N - 3N$ all nodes are at level $l + 1$.*

The Algorithm

- Complexity

- Time (under assumption of initial awakening it is $5N \log_2 N$)

*At the highest level only **Test**, **Reject**, and **Report** messages are used.*

The Algorithm

- Complexity
 - Time (under assumption of initial awakening it is $5N \log_2 N$)

As a result we have the algorithm complete under $5N \log_2 N$ time units.